

Supplemental: Introduction to Data Warehousing

ISM 6562 - Big Data for Business Applications

Dr. Tim Smith

2026-01-01

Table of contents

Why This Document Exists	2
Two Kinds of Databases	2
Operational Databases (OLTP)	2
Analytical Databases (OLAP)	4
Side-by-Side Comparison	4
Normalization vs Denormalization	5
Normalization (Operational Databases)	5
Denormalization (Analytical Databases)	6
The Star Schema	6
Fact Tables	7
Dimension Tables	7
How the Star Schema Looks	8
Why Stars Are Fast	9
Surrogate Keys	10
Other Common Warehouse Schemas	11
Snowflake Schema	11
Galaxy Schema (Multi-Fact)	11
The ETL Pipeline	11
Connecting to the Midterm	12

Key Takeaways	13
Further Reading	13

Why This Document Exists

The midterm project asks you to work with a **data warehouse** built using a **star schema**. If you haven't taken a dedicated data warehousing course, some of these concepts may be unfamiliar. This supplemental reading gives you enough background to understand what the midterm's warehouse is doing and why it's designed that way.

You do not need to memorize all of this. The goal is to understand the core ideas so that the midterm assignment makes sense.

Two Kinds of Databases

Most organizations maintain two fundamentally different types of databases, each optimized for a different purpose.

Operational Databases (OLTP)

An **operational database** (also called a transactional database or OLTP system) supports the day-to-day operations of a business. When a customer places an order, when an employee clocks in, when a product's inventory is updated — these events are recorded in operational databases.

Characteristics:

- **Purpose:** Record individual business transactions as they happen
- **Users:** Applications, front-end systems, employees entering data
- **Queries:** Simple reads and writes — insert one order, update one customer, look up one product
- **Optimized for:** Fast inserts and updates on individual rows
- **Schema:** Highly **normalized** (explained below) to avoid data duplication

In the midterm, the `sales-db` and `hr-db` are operational databases. They store customers, products, orders, employees, departments, and timesheets in normalized tables.

The sales operational database has three tables — `customers`, `products`, and `orders` — with foreign key relationships:

Sales Operational Database (sales-db) — Normalized

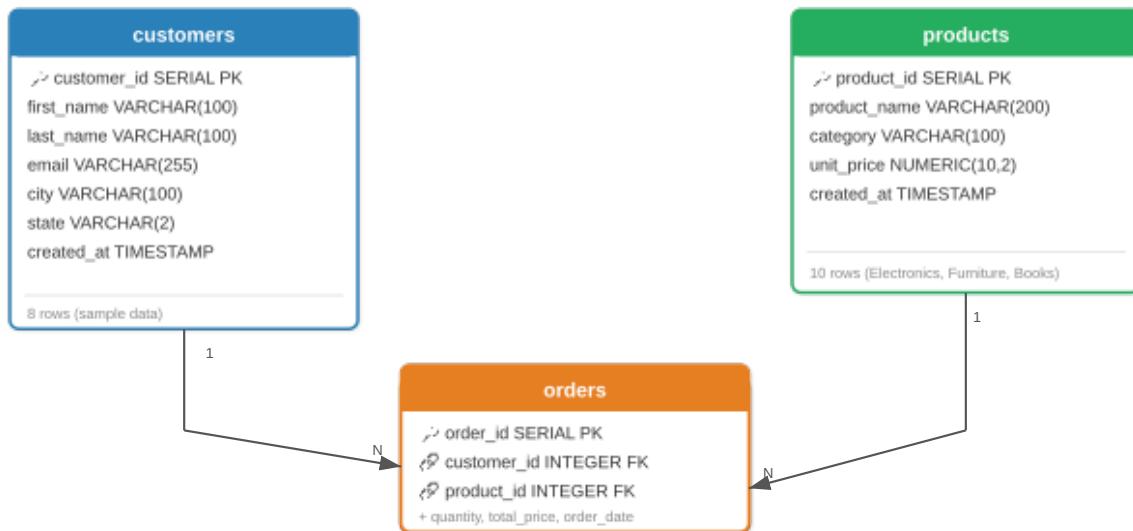


Figure 1: Sales Operational Database (sales-db) — Entity Relationship Diagram

The HR operational database has three tables — **departments**, **employees**, and **timesheets** — also with foreign key relationships:

HR Operational Database (hr-db) — Normalized



Normalized Design (3NF)

"Engineering" is stored once in departments. Employees reference it via department_id.
To get an employee's department name, you must JOIN the two tables.

Figure 2: HR Operational Database (hr-db) — Entity Relationship Diagram

Analytical Databases (OLAP)

An **analytical database** (also called a data warehouse or OLAP system) is designed for answering business questions across large amounts of data. Rather than recording individual transactions, it's built to support queries like "What was total revenue by product category last quarter?" or "Which department has the highest employee turnover?"

Characteristics:

- **Purpose:** Support analytical queries and business intelligence
- **Users:** Analysts, managers, reporting tools, dashboards
- **Queries:** Complex reads across many rows — aggregations, joins across multiple dimensions, trend analysis
- **Optimized for:** Fast reads over large datasets (not fast writes)
- **Schema:** **Denormalized** (explained below) to minimize the number of joins needed for analytical queries

In the midterm, the `warehouse-db` is the analytical database. It stores data in a star schema designed for fast analytical queries.

Side-by-Side Comparison

Aspect	Operational (OLTP)	Analytical (OLAP)
Primary purpose	Record transactions	Analyze trends
Typical query	"Get order #1234"	"Total revenue by quarter and region"
Data freshness	Real-time	Periodically refreshed (ETL)
Schema design	Normalized (3NF)	Denormalized (star/snowflake)
Number of tables	Many (to avoid duplication)	Fewer (to reduce joins)
Query complexity	Simple (1–2 table lookups)	Complex (joins across many dimensions)
Write pattern	Many small writes	Bulk loads (ETL batches)
Read pattern	Few rows at a time	Millions of rows aggregated

Normalization vs Denormalization

Normalization (Operational Databases)

Normalization is the process of organizing data to eliminate redundancy. In a normalized database, each piece of information is stored in exactly one place. If you need to change a department's name, you change it in one row of the `departments` table, and every employee in that department automatically reflects the change.

Example from the midterm's HR operational database:

`departments` table:

	department_id	department_name	location
1		Engineering	Tampa
2		Marketing	Orlando

`employees` table:

	employee_id	first_name	last_name	department_id
101		Alice	Johnson	1
102		Bob	Smith	1
103		Carol	Davis	2

Notice that "Engineering" and "Tampa" are stored only once (in the `departments` table). Employees reference their department via `department_id`. To get an employee's department name, you must JOIN the two tables:

```
SELECT e.first_name, e.last_name, d.department_name, d.location
FROM employees e
JOIN departments d ON e.department_id = d.department_id;
```

Benefits of normalization:

- No duplicated data — change a department name in one place
- Less storage waste
- Easier to maintain data consistency

Drawback: Analytical queries require many JOINs, which can be slow at scale.

Denormalization (Analytical Databases)

Denormalization is the deliberate introduction of redundancy to speed up read queries. In a denormalized table, related data is stored together so you don't need JOINs to answer common questions.

Example from the midterm's warehouse database:

dim_employee table:					
	employee_key	first_name	last_name	department_name	department_location
1		Alice	Johnson	Engineering	Tampa
2		Bob	Smith	Engineering	Tampa
3		Carol	Davis	Marketing	Orlando

Notice that “Engineering” and “Tampa” are repeated for both Alice and Bob. This is intentional — it means any query about employees and their departments can be answered from this single table without a JOIN:

```
SELECT first_name, last_name, department_name, department_location  
FROM dim_employee;
```

Benefits of denormalization:

- Faster queries (fewer JOINs)
- Simpler query logic
- Better performance for aggregations across large datasets

Drawback: Data is duplicated, so updates require changing multiple rows. This is acceptable in a warehouse because data is loaded in bulk by an ETL process, not updated row-by-row.

The Star Schema

The **star schema** is the most common design pattern for data warehouses. It gets its name from its shape: a central **fact table** surrounded by **dimension tables**, like a star.

Fact Tables

A **fact table** stores the measurable events of a business — the things you want to count, sum, or average. Each row represents one business event at a specific grain (level of detail).

Key characteristics:

- Contains **measures** (numeric values): quantity, price, revenue, duration
- Contains **foreign keys** pointing to dimension tables
- Typically the largest table in the warehouse (many rows)
- Each row represents one atomic business event

In the midterm, `fact_sales` is the fact table:

```
fact_sales:  
    sale_key          -- Surrogate primary key  
    date_key          -- FK → dim_date  
    customer_key      -- FK → dim_customer  
    product_key       -- FK → dim_product  
    source_order_id   -- Link back to operational database  
    quantity          -- Measure: how many units  
    unit_price        -- Measure: price per unit  
    total_price       -- Measure: total revenue
```

Each row represents one order. The measures (`quantity`, `unit_price`, `total_price`) are the numbers you aggregate in analytical queries.

Dimension Tables

A **dimension table** stores the descriptive attributes that give context to the facts. They answer the “who, what, when, where” of each business event.

Key characteristics:

- Contains **descriptive attributes** (text, categories, dates)
- Relatively few rows compared to fact tables
- Denormalized — related attributes are stored together
- Provide the labels and groupings for analytical queries

In the midterm, there are four dimension tables:

`dim_date` — When did the sale happen?

```

dim_date:
date_key      -- Integer key (YYYYMMDD format, e.g., 20260115)
full_date     -- The actual date
year, quarter, month, day_of_month
month_name    -- "January", "February", etc.
day_of_week, day_name, is_weekend

```

The date dimension is pre-populated with every day of the year. Pre-calculated fields like `quarter`, `month_name`, and `is_weekend` mean your queries never need date arithmetic — just filter or group by the column you need.

dim_customer — Who bought it?

```

dim_customer:
customer_key      -- Warehouse surrogate key
source_customer_id -- Link to operational sales-db
first_name, last_name, email, city, state

```

dim_product — What was bought?

```

dim_product:
product_key      -- Warehouse surrogate key
source_product_id -- Link to operational sales-db
product_name, category, unit_price

```

dim_employee — Who works here? (with denormalized department info)

```

dim_employee:
employee_key      -- Warehouse surrogate key
source_employee_id -- Link to operational hr-db
first_name, last_name, email, job_title
department_name    -- Denormalized from departments table
department_location -- Denormalized from departments table

```

How the Star Schema Looks

The diagram below shows the complete star schema used in the midterm's warehouse. The central `fact_sales` table (red) is surrounded by four dimension tables (blue, green, orange, purple). Foreign key relationships are shown as arrows; the dashed line to `dim_employee` indicates it is not directly referenced by `fact_sales` via a foreign key but is part of the warehouse design.



Figure 3: Data Warehouse (warehouse-db) — Star Schema ERD

The fact table sits at the center. Each dimension table connects to it through a foreign key. This is the “star” shape.

Why Stars Are Fast

Consider this business question: “*What was total revenue by product category and quarter in 2026?*”

In the warehouse (star schema) — one simple query:

```

SELECT p.category, d.quarter, SUM(f.total_price) AS revenue
FROM fact_sales f
JOIN dim_product p ON f.product_key = p.product_key
JOIN dim_date d ON f.date_key = d.date_key
WHERE d.year = 2026
    
```

```
GROUP BY p.category, d.quarter
ORDER BY d.quarter, revenue DESC;
```

In the operational database — more complex, slower:

```
SELECT p.category,
       EXTRACT(QUARTER FROM o.order_date) AS quarter,
       SUM(o.quantity * p.unit_price) AS revenue
  FROM orders o
 JOIN products p ON o.product_id = p.product_id
 WHERE EXTRACT(YEAR FROM o.order_date) = 2026
 GROUP BY p.category, EXTRACT(QUARTER FROM o.order_date)
 ORDER BY quarter, revenue DESC;
```

The warehouse version is simpler (pre-calculated quarter, pre-computed total_price) and faster at scale (fewer rows to scan, indexed foreign keys, denormalized dimensions).

Surrogate Keys

You may have noticed that the dimension tables use two kinds of keys:

- **Surrogate key** (`customer_key`, `product_key`, etc.) — an auto-generated integer that serves as the primary key in the warehouse
- **Source key** (`source_customer_id`, `source_product_id`, etc.) — the original primary key from the operational database

Why not just use the operational database's key?

Surrogate keys provide several benefits:

1. **Stability:** If the operational system changes its key format (e.g., migrates from integer IDs to UUIDs), the warehouse is unaffected
2. **History tracking:** You can keep multiple versions of a customer record (e.g., before and after an address change) with different surrogate keys
3. **Cross-system integration:** If you load data from multiple operational systems, their keys might conflict; surrogate keys avoid collisions
4. **Performance:** Integer surrogate keys are compact and fast for joins

Other Common Warehouse Schemas

The star schema used in the midterm is the most common pattern, but you should be aware of two alternatives.

Snowflake Schema

A **snowflake schema** is a star schema where some dimension tables are further normalized into sub-tables. For example, instead of storing `department_name` directly in `dim_employee`, you might have a separate `dim_department` table that `dim_employee` references.

```
dim_department    dim_employee    fact_sales    dim_product    dim_category  
                    dim_date
```

Pros: Less data redundancy in dimensions; may save storage.

Cons: More JOINs needed for queries; more complex to understand and maintain.

In practice, star schemas are more common than snowflake schemas because the small amount of redundancy in dimension tables is a worthwhile trade-off for simpler, faster queries.

Galaxy Schema (Multi-Fact)

A **galaxy schema** (also called a fact constellation) has multiple fact tables that share some dimension tables. For example, a business might have both a `fact_sales` table and a `fact_returns` table, both referencing `dim_customer`, `dim_product`, and `dim_date`.

This is common in real-world warehouses but is beyond the scope of the midterm.

The ETL Pipeline

Data doesn't magically appear in the warehouse. An **ETL pipeline** moves data from operational databases to the warehouse:

1. **Extract:** Read data from operational databases (sales-db, hr-db)
2. **Transform:** Clean, reshape, and denormalize the data (e.g., join employees with departments, generate date dimension rows, compute derived fields)
3. **Load:** Insert the transformed data into the warehouse tables

The following diagram shows how the midterm's ETL pipeline moves data from the two operational databases into the star-schema warehouse:

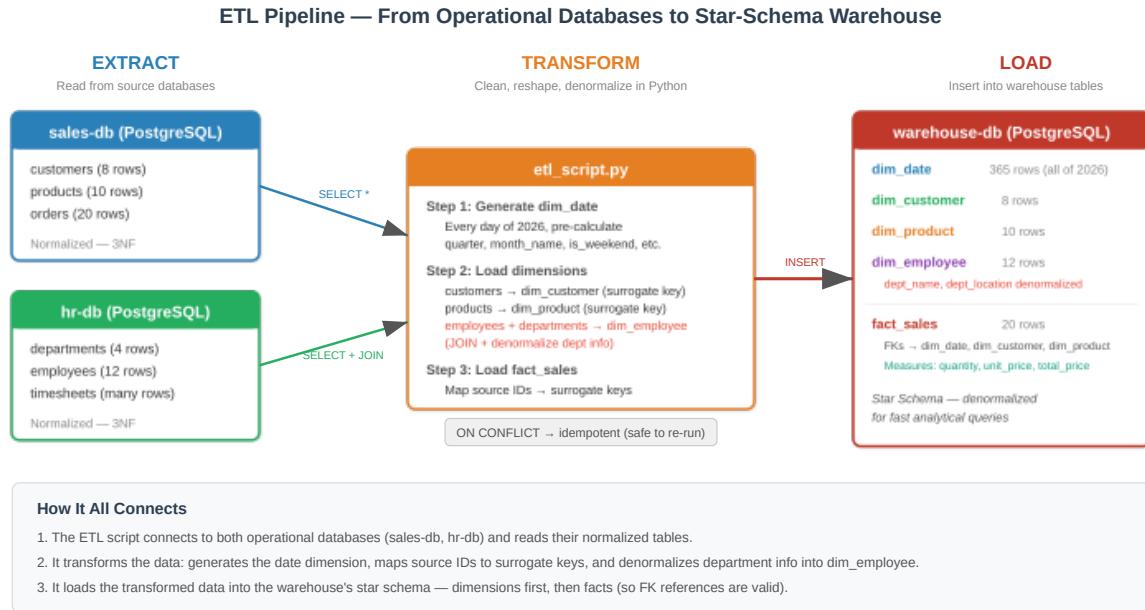


Figure 4: ETL Pipeline — From Operational Databases to Star-Schema Warehouse

In the midterm, the `etl_script.py` performs all three steps. Key design decisions in the midterm's ETL:

- **Idempotent loading:** The script uses `ON CONFLICT` clauses so it can be run multiple times without duplicating data. If a customer already exists, it updates the record; if an order already exists, it skips it.
- **Dimension loading before facts:** Dimensions are loaded first so that the foreign key references in `fact_sales` are valid.
- **Date dimension generation:** The date dimension is generated programmatically (every day of 2026) rather than extracted from an operational system.

Connecting to the Midterm

Here's how all of these concepts map to what you'll see in the midterm project:

Concept	Where You'll See It
Operational databases (OLTP)	<code>sales-db</code> and <code>hr-db</code> — normalized tables with JOINs
Analytical database (OLAP)	<code>warehouse-db</code> — star schema with denormalized dimensions
Fact table	<code>fact_sales</code> — one row per order, with measures (quantity, price)
Dimension tables	<code>dim_date</code> , <code>dim_customer</code> , <code>dim_product</code> , <code>dim_employee</code>
Denormalization	<code>dim_employee</code> includes <code>department_name</code> and <code>department_location</code> directly
Surrogate keys	<code>customer_key</code> , <code>product_key</code> , etc. (auto-generated, separate from source IDs)
Pre-calculated attributes	<code>dim_date</code> has <code>quarter</code> , <code>month_name</code> , <code>is_weekend</code> already computed
ETL pipeline	<code>etl_script.py</code> extracts from operational DBs, transforms, loads into warehouse
Idempotent loading	<code>ON CONFLICT</code> clauses in the load step prevent duplicate data

Key Takeaways

1. **Operational databases** are for recording transactions; **warehouses** are for analyzing them
2. **Normalization** (operational) eliminates redundancy; **denormalization** (warehouse) speeds up queries
3. A **star schema** has a central fact table (measures) surrounded by dimension tables (descriptive context)
4. **Fact tables** store what happened (quantities, amounts); **dimension tables** describe who, what, when, where
5. **Surrogate keys** decouple the warehouse from operational system changes
6. **ETL pipelines** move and reshape data from operational systems into the warehouse
7. The midterm uses all of these patterns — the supplemental readings and the assignment will walk you through each one

Further Reading

If you want to go deeper into data warehousing, these are good starting points:

- Ralph Kimball, *The Data Warehouse Toolkit* (the definitive reference for star schema design)

- [Kimball Group Design Tips](#) — short articles on specific design patterns
- [Star Schema vs Snowflake Schema](#) — IBM overview of schema patterns